

DSC 630 - Predictive Analytics

Term Project - Summer 2024

David Berberena | Brian Mann

```
In [1]: # import libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression, Ridge
from sklearn.metrics import r2_score, mean_squared_error
from sklearn.metrics import accuracy_score, precision_recall_fscore_support
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor, StackingRegressor
from sklearn.base import BaseEstimator, RegressorMixin
import statsmodels.api as sm
```

```
In [2]: # hide warnings
import warnings
warnings.filterwarnings('ignore')
from sklearn.exceptions import ConvergenceWarning
```

Milestone 3 - Preliminary Analysis

```
In [3]: # dataset
url = 'https://www.kaggle.com/datasets/mrsimple07/restaurants-revenue-prediction'
```

Step 1: Reading and Describing the Data

```
In [4]: # read in the data and show the first few rows
df = pd.read_csv('restaurant_revenue.csv')
df.head()
```

```
Out[4]:
```

	Number_of_Customers	Menu_Price	Marketing_Spend	Cuisine_Type	Average_Customer_Spending	Promotions	Reviews	Monthly_R
0	61	43.117635	12.663793	Japanese	36.236133	0	45	350.
1	24	40.020077	4.577892	Italian	17.952562	0	36	221.
2	81	41.981485	4.652911	Japanese	22.600420	1	91	326.
3	70	43.005307	4.416053	Italian	18.984098	1	59	348.
4	30	17.456199	3.475052	Italian	12.766143	1	30	185.

```
In [5]: # get the size and shape
df.shape
```

```
Out[5]: (1000, 8)
```

```
In [6]: # check for missing values
df.isna().sum()
```

```
Out[6]: Number_of_Customers    0
Menu_Price                    0
Marketing_Spend                0
Cuisine_Type                  0
Average_Customer_Spending     0
Promotions                    0
Reviews                       0
Monthly_Revenue               0
dtype: int64
```

```
In [7]: # assess common stats for numerical variables
df.describe()
```

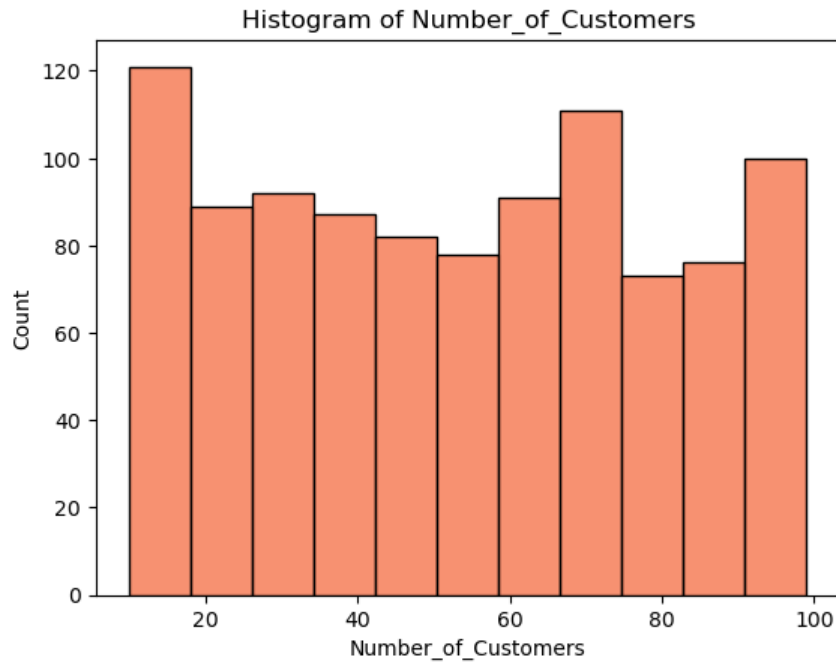
```
Out[7]:
```

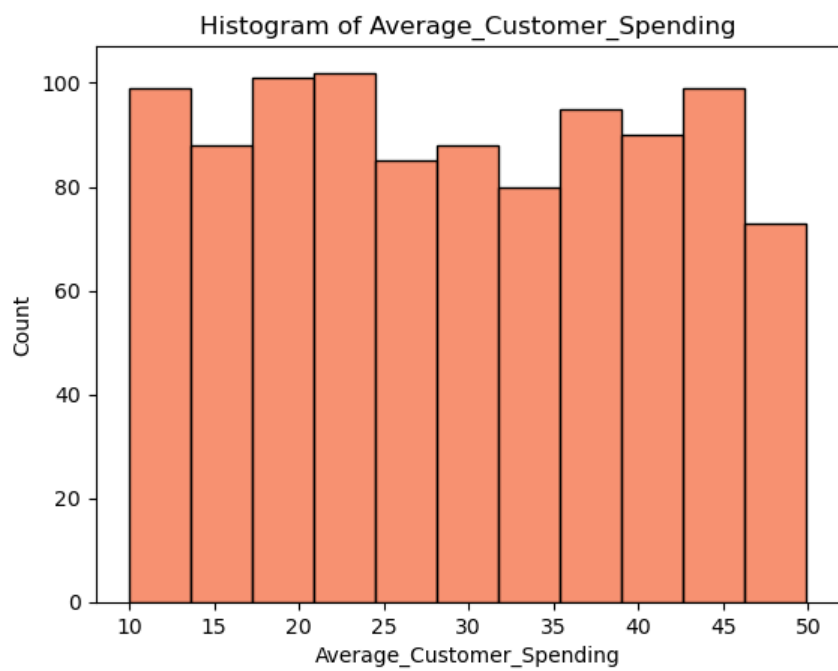

	Number_of_Customers	Menu_Price	Marketing_Spend	Average_Customer_Spending	Promotions	Reviews	Monthly_Revenue
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000
mean	53.271000	30.219120	9.958726	29.477085	0.497000	49.837000	268.72417
std	26.364914	11.278760	5.845586	11.471686	0.500241	29.226334	103.98295
min	10.000000	10.009501	0.003768	10.037177	0.000000	0.000000	-28.97780
25%	30.000000	20.396828	4.690724	19.603041	0.000000	24.000000	197.10364
50%	54.000000	30.860614	10.092047	29.251365	0.000000	50.000000	270.21396
75%	74.000000	39.843868	14.992436	39.553220	1.000000	76.000000	343.39579
max	99.000000	49.974140	19.994276	49.900725	1.000000	99.000000	563.38133


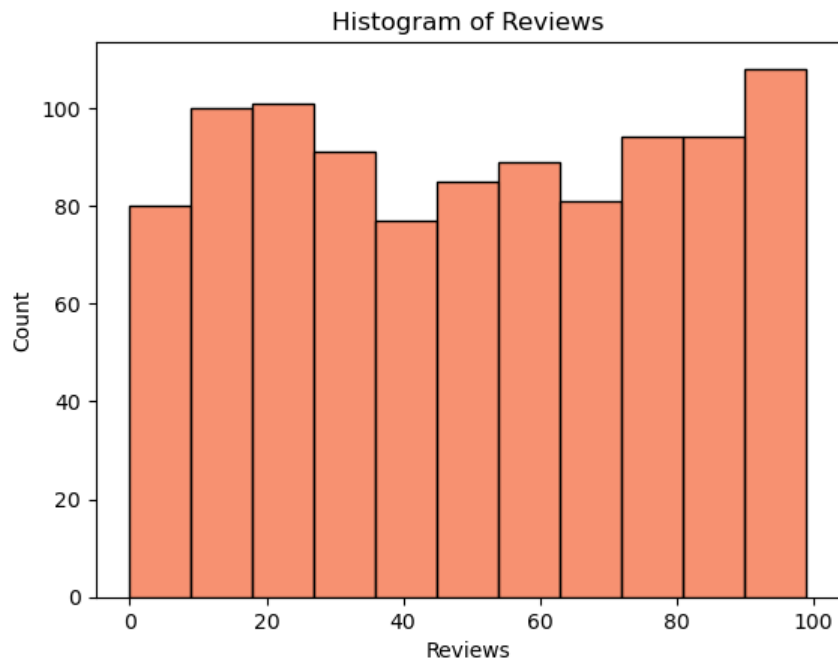
Step 2: Histograms and Bar Charts

```
In [8]: # this function takes in a df and a column and prints out a histogram
def make_histogram(data, col):
    sns.histplot(data[col], color='#f56c42')
    plt.title(f'Histogram of {col}')
    plt.show()
```

```
In [9]: # print out histograms for each of the primary numerical variables
for col in df.describe().drop(columns='Promotions').columns:
    make_histogram(df, col)
```

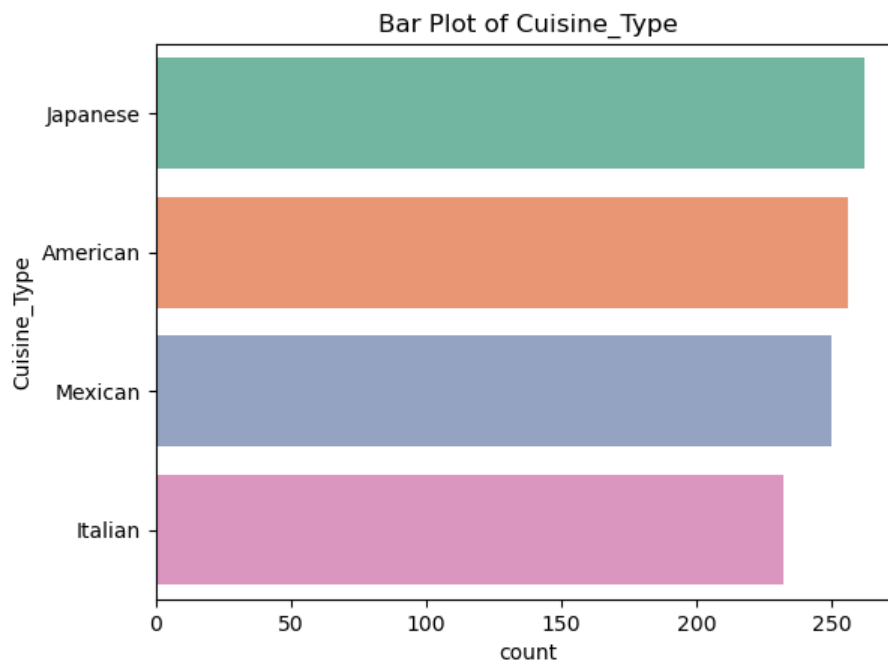
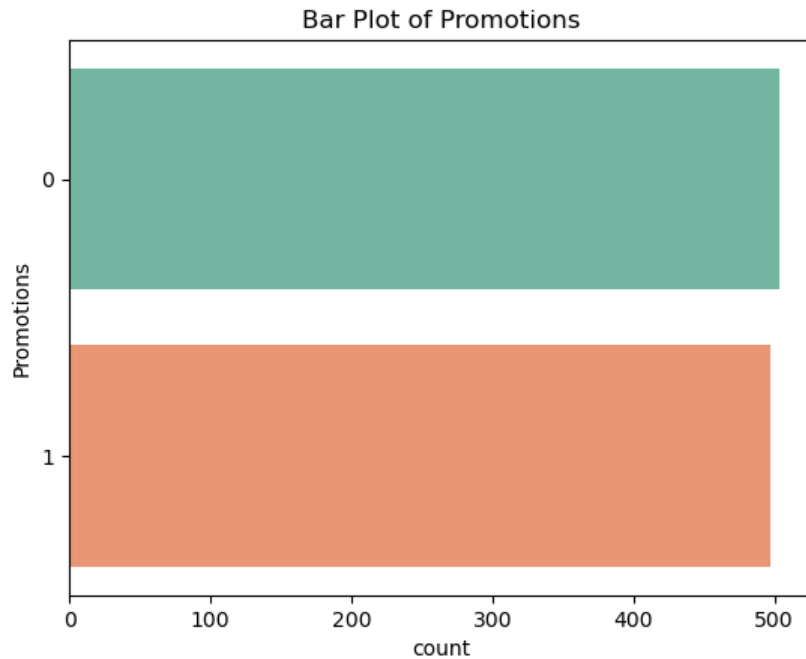






```
In [10]: # this function takes in a df and a column and prints out a horizontal bar
# plot of the counts
def count_bar(data, col):
    counts = data[col].value_counts().reset_index()
    sns.barplot(data=counts, y=col, x='count', orient='h', palette='Set2')
    plt.title(f'Bar Plot of {col}')
    plt.show()
```

```
In [11]: # print out bar plots for the counts of promotions and cuisine type
for col in ['Promotions', 'Cuisine_Type']:
    count_bar(df, col)
```



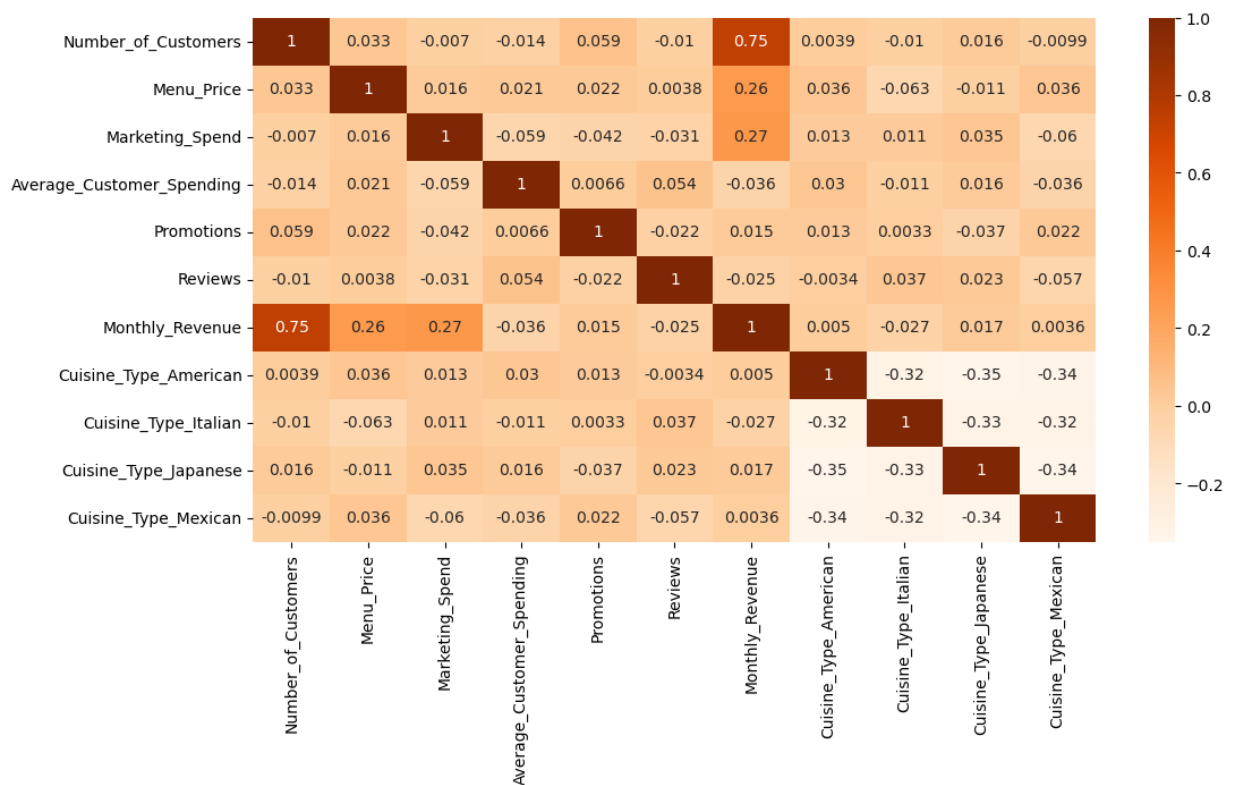
Step 3: Correlation and Scatterplots

```
In [12]: # create dummy variables for cuisine type and make it its own df
dum = pd.get_dummies(df, columns=['Cuisine_Type'])
dum.head()
```

```
Out[12]:
```


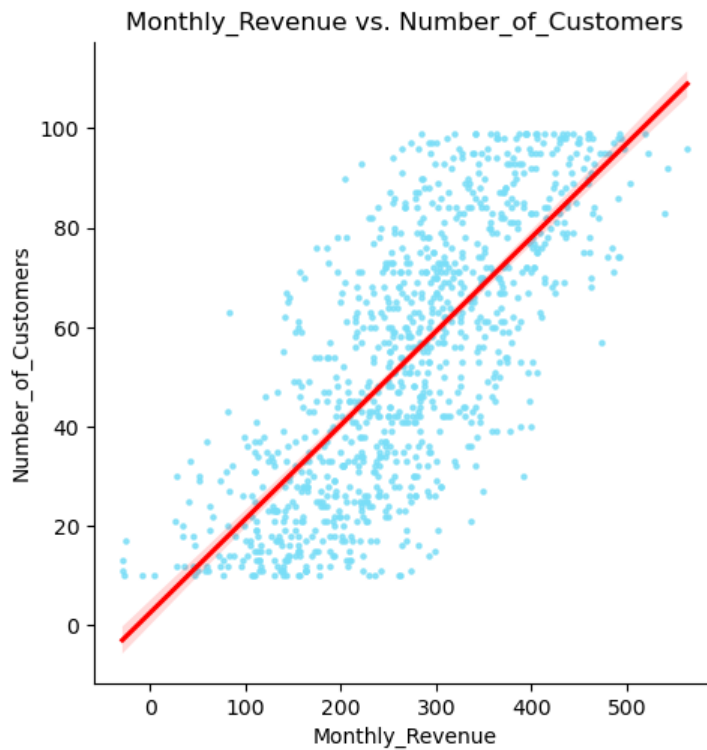
	Number_of_Customers	Menu_Price	Marketing_Spend	Average_Customer_Spending	Promotions	Reviews	Monthly_Revenue	Cuisin
0	61	43.117635	12.663793	36.236133	0	45	350.912040	
1	24	40.020077	4.577892	17.952562	0	36	221.319091	
2	81	41.981485	4.652911	22.600420	1	91	326.529763	
3	70	43.005307	4.416053	18.984098	1	59	348.190573	
4	30	17.456199	3.475052	12.766143	1	30	185.009121	

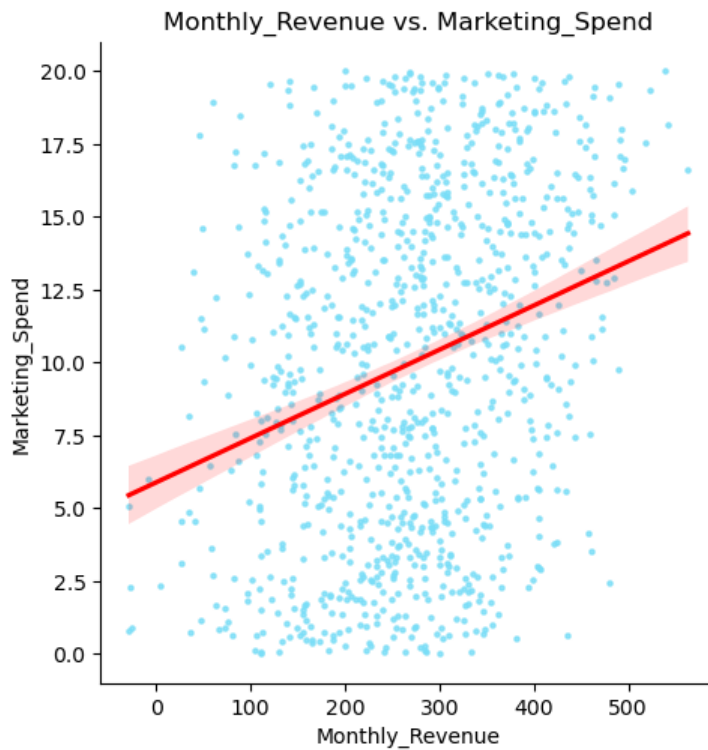
```
In [13]: # find the correlation among all of the variables in the df
corrs = dum.corr()
plt.figure(figsize=(12,6))
# plot a heatmap of the correlations
sns.heatmap(corrs, cmap="Oranges", annot=True)
plt.show()
```



```
In [14]: # this function takes in a df and two columns, then prints out a scatter plot
# with a line of best fit containing margins of error
def scatter_with_fit(data, col1, col2):
    sns.lmplot(data=data, x=col1, y=col2, line_kws={'color': 'red'},
               scatter_kws={'color': '#79def7', 's': 5})
    plt.title(f'{col1} vs. {col2}')
    plt.show()
```

```
In [15]: # print out scatterplots with best fit lines for the variables with
# high correlations with monthly revenue
for col in ['Number_of_Customers', 'Menu_Price', 'Marketing_Spend']:
    scatter_with_fit(df, 'Monthly_Revenue', col)
```






Step 4: Grouping Revenue by Categories

```
In [16]: # group cuisine types and promotions by average monthly revenue
types = df.groupby(by='Cuisine_Type').Monthly_Revenue.agg(np.mean).reset_index()
proms = df.groupby(by='Promotions').Monthly_Revenue.agg(np.mean).reset_index()
```

```
In [17]: # print a barplot of revenue by cuisine type
sns.barplot(data=types, x='Monthly_Revenue', y='Cuisine_Type')
plt.show()
```




```
In [18]: types
```

```
Out[18]:
```

	Cuisine_Type	Monthly_Revenue
0	American	269.604825
1	Italian	263.699862
2	Japanese	271.700186
3	Mexican	269.366082

```
In [19]: # rename the 0's to NO and the 1's to YES
proms.iloc[0,0] = 'NO'
proms.iloc[1,0] = 'YES'
# print a barplot of revenue by whether or not there was a promotion
sns.barplot(data=proms, x='Monthly_Revenue', y='Promotions')
plt.show()
```



```
In [20]: proms
```

```
Out[20]:
```

	Promotions	Monthly_Revenue
0	NO	267.188084
1	YES	270.278805

```
In [21]: # Percent difference in aggregated monthly revenue for the presence of a promotional campaign
# to establish the campaign's worth

print('%diff in aggregated monthly restaurant revenue with ' +
      'the presence of a promotional campaign: ',
      ((proms['Monthly_Revenue'][1] / proms['Monthly_Revenue'][0]) - 1) * 100)

%diff in aggregated monthly restaurant revenue with the presence of a promotional campaign:  1.1567584291771071
```

Step 5: Modeling

```
In [22]: # get the dummy variables, while also dropping one of them
dum = pd.get_dummies(df, columns=['Cuisine_Type'], drop_first=True)
# set the features to be all of the variables except monthly revenue
features = dum.drop(columns='Monthly_Revenue')
# split the data into training and testing sets at an 80/20 ratio
xtrain, xtest, ytrain, ytest = train_test_split(features,
                                                dum['Monthly_Revenue'], test_size=0.2, random_state=42)
```

```

In [32]: # Create a custom wrapper for statsmodels OLS
class OLSWrapper(BaseEstimator, RegressorMixin):
    def __init__(self):
        self.model = None

    def fit(self, X, y):
        # Add a constant to the input features to account for the intercept
        X = sm.add_constant(X)
        self.model = sm.OLS(y, X).fit()
        print(self.model.summary())
        return self

    def predict(self, X):
        # Add a constant for prediction
        X = sm.add_constant(X)
        return self.model.predict(X)

    def score(self, X, y):
        predictions = self.predict(X)
        # Return R-squared score
        return 1 - np.sum((y - predictions) ** 2) / np.sum((y - np.mean(y)) ** 2)

# Create a pipeline with a custom wrapper
pipe = Pipeline([
    ('scale', StandardScaler()),
    ('model', OLSWrapper())
])

# Set up a search space
search_space = [{'model': [OLSWrapper()]},
                 {'model': [Ridge(max_iter=10000)],
                  'model__alpha': np.logspace(-4, 4, 50)}]

# Perform grid search with cross-validation
grid_search = GridSearchCV(pipe, search_space, cv=5, scoring='r2')
gs = grid_search.fit(xtrain, ytrain)

be = gs.best_estimator_
print(f'Best estimator: {be}')
# make a prediction using the test data
pred = be.predict(xtest)
# get R^2 and RMSE
r2 = r2_score(ytest, pred)
rmse = np.sqrt(mean_squared_error(ytest, pred))
print(f'R^2: {r2}')
print(f'RMSE: {rmse}')

```

OLS Regression Results

```

=====
Dep. Variable:    Monthly_Revenue    R-squared:                0.680
Model:            OLS                Adj. R-squared:           0.675
Method:           Least Squares      F-statistic:              148.7
Date:             Fri, 26 Jul 2024   Prob (F-statistic):       1.83e-149
Time:             02:22:11           Log-Likelihood:           -3505.1
No. Observations: 640               AIC:                      7030.
Df Residuals:     630               BIC:                      7075.
Df Model:         9
Covariance Type:  nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	268.9217	2.305	116.692	0.000	264.396	273.447
x1	75.8256	2.316	32.734	0.000	71.277	80.374
x2	23.7732	2.320	10.249	0.000	19.218	28.328
x3	26.8290	2.318	11.574	0.000	22.277	31.381
x4	-1.6173	2.317	-0.698	0.485	-6.167	2.932
x5	-4.1567	2.319	-1.793	0.074	-8.710	0.397
x6	-1.0656	2.316	-0.460	0.646	-5.614	3.483
x7	0.2211	2.752	0.080	0.936	-5.183	5.625
x8	-1.9680	2.783	-0.707	0.480	-7.432	3.496
x9	0.9550	2.782	0.343	0.731	-4.507	6.417
=====						
Omnibus:		2.109	Durbin-Watson:			1.858
Prob(Omnibus):		0.348	Jarque-Bera (JB):			2.068
Skew:		-0.139	Prob(JB):			0.356
Kurtosis:		2.995	Cond. No.			1.98
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

OLS Regression Results

```

=====
Dep. Variable:    Monthly_Revenue    R-squared:                0.681
Model:            OLS                Adj. R-squared:           0.676
Method:           Least Squares      F-statistic:              149.3
Date:             Fri, 26 Jul 2024   Prob (F-statistic):       7.29e-150
Time:             02:22:11           Log-Likelihood:           -3497.4
No. Observations: 640               AIC:                      7015.
Df Residuals:     630               BIC:                      7059.
Df Model:         9
Covariance Type:  nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	268.7920	2.277	118.048	0.000	264.321	273.263
x1	75.2285	2.288	32.878	0.000	70.735	79.722
x2	26.1852	2.296	11.403	0.000	21.676	30.695
x3	25.7532	2.291	11.241	0.000	21.254	30.252
x4	0.4491	2.281	0.197	0.844	-4.030	4.928
x5	-3.6046	2.297	-1.569	0.117	-8.116	0.907
x6	1.5748	2.293	0.687	0.492	-2.928	6.078
x7	-0.1051	2.846	-0.037	0.971	-5.695	5.485
x8	-2.0513	2.888	-0.710	0.478	-7.722	3.619
x9	-0.9865	2.875	-0.343	0.732	-6.633	4.660
=====						
Omnibus:		0.811	Durbin-Watson:		1.889	
Prob(Omnibus):		0.667	Jarque-Bera (JB):		0.877	
Skew:		-0.013	Prob(JB):		0.645	
Kurtosis:		2.821	Cond. No.		2.16	
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

OLS Regression Results

```

=====
Dep. Variable:    Monthly_Revenue    R-squared:                0.692
Model:            OLS                Adj. R-squared:           0.687
Method:           Least Squares      F-statistic:              156.9
Date:             Fri, 26 Jul 2024   Prob (F-statistic):       1.76e-154
Time:             02:22:11           Log-Likelihood:           -3501.0
No. Observations: 640               AIC:                      7022.

```

Df Residuals: 630 BIC: 7067.
Df Model: 9
Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
const	272.0739	2.290	118.819	0.000	267.577	276.571
x1	75.8445	2.301	32.956	0.000	71.325	80.364
x2	25.5013	2.300	11.086	0.000	20.984	30.018
x3	27.4173	2.312	11.861	0.000	22.878	31.957
x4	-0.8585	2.298	-0.374	0.709	-5.371	3.654
x5	-3.6514	2.302	-1.586	0.113	-8.171	0.869
x6	-1.0937	2.301	-0.475	0.635	-5.612	3.424
x7	-0.9335	2.765	-0.338	0.736	-6.364	4.497
x8	-2.6698	2.808	-0.951	0.342	-8.184	2.844
x9	-0.9220	2.802	-0.329	0.742	-6.425	4.581

Omnibus: 2.153 Durbin-Watson: 1.816
Prob(Omnibus): 0.341 Jarque-Bera (JB): 2.218
Skew: -0.137 Prob(JB): 0.330
Kurtosis: 2.909 Cond. No. 2.04

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
OLS Regression Results

Dep. Variable: Monthly_Revenue R-squared: 0.673
Model: OLS Adj. R-squared: 0.669
Method: Least Squares F-statistic: 144.2
Date: Fri, 26 Jul 2024 Prob (F-statistic): 1.17e-146
Time: 02:22:11 Log-Likelihood: -3511.7
No. Observations: 640 AIC: 7043.
Df Residuals: 630 BIC: 7088.
Df Model: 9
Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
const	270.4416	2.329	116.139	0.000	265.869	275.014
x1	74.6859	2.338	31.950	0.000	70.095	79.276
x2	24.5162	2.336	10.496	0.000	19.930	29.103
x3	26.7485	2.350	11.383	0.000	22.134	31.363
x4	0.4820	2.344	0.206	0.837	-4.122	5.086
x5	-3.7136	2.353	-1.578	0.115	-8.334	0.907
x6	-2.4734	2.344	-1.055	0.292	-7.076	2.129
x7	2.4353	2.819	0.864	0.388	-3.101	7.972
x8	1.5313	2.854	0.537	0.592	-4.073	7.136
x9	1.4608	2.862	0.510	0.610	-4.160	7.082

Omnibus: 0.354 Durbin-Watson: 1.849
Prob(Omnibus): 0.838 Jarque-Bera (JB): 0.388
Skew: -0.056 Prob(JB): 0.824
Kurtosis: 2.957 Cond. No. 2.05

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
OLS Regression Results

Dep. Variable: Monthly_Revenue R-squared: 0.685
Model: OLS Adj. R-squared: 0.681
Method: Least Squares F-statistic: 152.5
Date: Fri, 26 Jul 2024 Prob (F-statistic): 8.41e-152
Time: 02:22:11 Log-Likelihood: -3501.1
No. Observations: 640 AIC: 7022.
Df Residuals: 630 BIC: 7067.
Df Model: 9
Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
const	269.2980	2.290	117.580	0.000	264.800	273.796
x1	75.1271	2.309	32.531	0.000	70.592	79.662
x2	23.8531	2.306	10.342	0.000	19.324	28.382

x3	24.6782	2.316	10.656	0.000	20.130	29.226
x4	-1.0789	2.303	-0.468	0.640	-5.601	3.444
x5	-3.1992	2.319	-1.379	0.168	-7.754	1.355
x6	-3.7074	2.302	-1.611	0.108	-8.228	0.813
x7	0.8092	2.792	0.290	0.772	-4.675	6.293
x8	1.6568	2.814	0.589	0.556	-3.870	7.184
x9	0.9453	2.796	0.338	0.735	-4.545	6.435

Omnibus:	0.884	Durbin-Watson:	1.863
Prob(Omnibus):	0.643	Jarque-Bera (JB):	0.883
Skew:	-0.090	Prob(JB):	0.643
Kurtosis:	2.976	Cond. No.	2.06

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Best estimator: Pipeline(steps=[('scale', StandardScaler()),
('model', Ridge(alpha=7.9060432109076855, max_iter=10000))])
R²: 0.6937014067272136
RMSE: 60.856889064989275

```
In [23]: # create a pipeline for a scale and model
pipe = Pipeline([('scale', StandardScaler()), ('model', LinearRegression())])

In [24]: # set up a search space for linear regression and ridge regression
search_space = [{'model': [LinearRegression()]},
                 {'model': [Ridge(max_iter=10000)],
                  'model__alpha': np.logspace(-4, 4, 50)}]

In [25]: # make a grid search cross-validation with 5-folds over the pipe and search space
grid_search = GridSearchCV(pipe, search_space, cv=5, n_jobs=-1, verbose=0, scoring='neg_mean_squared_error')

In [26]: # this function takes in training and testing variables, fits
# them to the grid search model, and prints out various results
def run_tests(x, xt, y, yt):
    # fit the training data to the grid search
    gs = grid_search.fit(x, y)
    # get the best estimator
    be = gs.best_estimator_
    print(f'Best estimator: {be}')
    # make a prediction using the test data
    pred = be.predict(xt)
    # get R^2 and RMSE
    r2 = r2_score(yt, pred)
    rmse = np.sqrt(mean_squared_error(yt, pred))
    print(f"RMSE: {rmse}")
    print(f"R^2: {r2}")
    # print out the list of variables with their respective coefficients
    coefs = be.named_steps['model'].coef_
    print("\nFeature importance:")
    for feature, coef in zip(features, coefs):
        print(f"{feature}: {coef:.4f}")

In [27]: run_tests(xtrain, xtest, ytrain, ytest)
```

Best estimator: Pipeline(steps=[('scale', StandardScaler()),
('model', Ridge(alpha=7.9060432109076855, max_iter=10000))])
RMSE: 60.856889064989275
R²: 0.6937014067272136

Feature importance:
Number_of_Customers: 74.5995
Menu_Price: 24.5487
Marketing_Spend: 26.0052
Average_Customer_Spending: -0.5611
Promotions: -3.5897
Reviews: -1.3689
Cuisine_Type_Italian: 0.4578
Cuisine_Type_Japanese: -0.6830
Cuisine_Type_Mexican: 0.2322

```
In [28]: # this function runs a tree regression model on the data using a parameter grid
def run_tree_regression():
    # set the initial state of the regressor
    tree_reg = DecisionTreeRegressor(random_state=1)
    # define a grid of potential parameters
    param_grid = {
        'max_depth': [None, 10, 20, 30, 40, 50],
        'min_samples_split': [2, 5, 10],
        'min_samples_leaf': [1, 2, 4]
    }
    # create a grid search using 5 cross-folds on the tree regression model
    grid_search = GridSearchCV(estimator=tree_reg, param_grid=param_grid,
                               cv=5, scoring='neg_mean_squared_error', n_jobs=-1, verbose=0)
    # apply the grid search to the training data
    grid_search.fit(xtrain, ytrain)
    # determine the best estimator
    be = grid_search.best_estimator_
    # use the best model to make predictions on the test data
    preds = be.predict(xtest)
    RMSE = np.sqrt(mean_squared_error(ytest, preds))
    R_squared = r2_score(ytest, preds)
    # print out the results of the best model's performance
    print(f'Best estimator: {be}')
    print(f'RMSE: {RMSE}')
    print(f'R^2: {R_squared}')
```

```
In [29]: run_tree_regression()
```

```
Best estimator: DecisionTreeRegressor(max_depth=10, min_samples_leaf=4, min_samples_split=10,
                                     random_state=1)
RMSE: 77.33858179980125
R^2: 0.5053273330434154
```

```
In [33]: # this function creates an ensemble model of ridge, tree, random forest and gradient boosting
# that applies a stacking regressor to the training data
def run_ensemble():
    # initialize each of the regression models
    ridge_reg = Ridge()
    tree_reg = DecisionTreeRegressor(random_state=1)
    rf_reg = RandomForestRegressor(random_state=1)
    gb_reg = GradientBoostingRegressor(random_state=1)
    # define the estimators for each model
    estimators = [
        ('ridge', ridge_reg),
        ('tree', tree_reg),
        ('rf', rf_reg),
        ('gb', gb_reg)
    ]
    # create a stack of models with a 5-fold cross-validation
    stacking_reg = StackingRegressor(
        estimators=estimators,
        cv=5
    )
    # fit the stack to the training data
    stacking_reg.fit(xtrain, ytrain)
    # use the fitted model to make predictions on the test data
    ypred_train = stacking_reg.predict(xtrain)
    ypred_test = stacking_reg.predict(xtest)
    # gather the assessment metrics
    rmse = np.sqrt(mean_squared_error(ytest, ypred_test))
    rsquared = r2_score(ytest, ypred_test)
    # print out the metrics
    print(f'RMSE: {rmse}')
    print(f'R^2: {rsquared}')
```

```
In [34]: run_ensemble()
```

```
RMSE: 61.785734154105945
R^2: 0.6842801198816617
```